

Investigating the isometric properties of Neural Machine Translation models on binary semantic-equivalence spaces

Atreya Shankar

Cognitive Systems, University of Potsdam

Department of Computational Linguistics, University of Zürich

atreya.shankar@{uni-potsdam.de, uzh.ch}

Abstract

Isometry is defined mathematically as a distance-preserving transformation between two metric spaces. In this research, we hypothesize that well-performing Neural Machine Translation (NMT) models function approximately isometrically on semantic metric spaces. That is to say, if two sentences are semantically equivalent on the source side, they should remain semantically equivalent after translation on the target side. We begin by utilizing two NMT models of varying performance to translate semantically-equivalent paraphrases based off WMT19 test data references. In order to quantify and simplify the notion of a semantic metric space, we treat it as a probabilistic binary semantic-equivalence space indicating either semantic equality or inequality; achieved by fine-tuning three transformer-based language models on Google’s PAWS-X paraphrase detection task. By using the paraphrase detection outputs, we investigate the frequency and composition of semantically isometric behaviour in the NMT models’ inputs and outputs.

1 Introduction

Isometry is defined mathematically as a distance-preserving transformation between two metric spaces (Coxeter, 1961). In this research, we view Neural Machine Translation (NMT) models from the perspective of semantic isometry and hypothesize that well-performing NMT models function approximately isometrically on semantic metric spaces. That is to say, if two sentences are semantically equivalent on the source side, they should remain semantically equivalent after translation on the target side given a well-performing NMT model. A simplified illustration of isometry in higher dimensional functional spaces can be seen in Figure 1.




Figure 1: Illustration of isometry in higher dimensional functional transformations (Hegde et al., 2015)

In order to conduct our investigation, we start by acquiring semantically equivalent paraphrases of WMT19 legacy and additional test references from Freitag et al. (2020) for en→de. Next, we utilize two NMT models of varying performance, specifically the SOTA FAIR’s WMT19 winning single transformer (Ng et al., 2019) and the non-SOTA Scaling NMT WMT16 Transformer (Ott et al., 2018), in order to translate the aforementioned paraphrases in the de→en translation direction. We use the former model pre-trained from fairseq (Ott et al., 2019) and train the latter model from scratch.

Next, we utilize three well-performing paraphrase detection models to approximate isometry in the NMT models’ translations. These paraphrase detection models are based off mBERT_{Base} (Devlin et al., 2019), XLM-R_{Base} (Conneau et al., 2019) and XLM-R_{Large} (Conneau et al., 2019) pre-trained multilingual language models; which are correspondingly fine-tuned on Google’s PAWS-X paraphrase detection task (Yang et al., 2019; Hu et al., 2020).

Using the outputs of the paraphrase detection models, we finally investigate the frequency and composition of semantically isometric behaviour in the NMT models’ inputs and outputs. We release our latest models and source code in our public GitHub repository¹.

¹<https://github.com/atreyasha/semantic-isometry-nmt>

2 Isometry and approximations

The concept of isometry in the context of semantic metric spaces can be *exactly* expressed as follows; where $s_i \in \mathbb{R}^{V \times N}$ refers to an input sentence’s tokenized matrix-form for vocabulary size V and maximum sentence length N , $f : \mathbb{R}^{V \times N} \rightarrow \mathbb{R}^{V' \times N'}$ refers to the NMT model’s inference function and $D_L : \mathbb{R}^{V \times 2N} \rightarrow \mathbb{R}_+$ refers to a semantic distance metric function for language L corresponding to the language of the respective sentences:

$$D_X(s_1, s_2) = D_Y(f(s_1), f(s_2)) \quad (1)$$

While elegant, this representation of isometry and a semantic distance metric is problematic. Firstly, exact isometry might not be a practical condition to achieve given real-life data instances with stochastic noise. Next, constructing continuous semantic metric spaces from discrete textual data is a difficult task and is in itself a developing field of research (Cer et al., 2017).

Approximate isometry: To address the first issue, we loosen the constraints of exact isometry to *approximate* isometry:

$$D_X(s_1, s_2) \approx D_Y(f(s_1), f(s_2)) \quad (2)$$

With this approximation, we can simplify the isometric relationship further into a binary semantic-equivalence function $S_L : \mathbb{R}^{V \times 2N} \rightarrow \{0, 1\}$, which compresses semantic distance metrics to semantic equality ($S_L = 1$) or inequality ($S_L = 0$) depending on some variable threshold $\delta_L \in \mathbb{R}_+$:

$$S_L(s_1, s_2) = \begin{cases} 1, & D_L(s_1, s_2) \leq \delta_L \\ 0, & D_L(s_1, s_2) > \delta_L \end{cases} \quad (3)$$

It is worth noting that the formulation in S_L is more meaningful for inferring isometry from semantic equality than from semantic inequality, due to the presence of a tighter bound for the former than the latter.

Probabilistic semantic-equivalence spaces: To address the second issue, we effectively delegate away the actual computation of a semantic distance metric and convert this into a probabilistic process; with a new definition for S_L below given a probability threshold ϵ with a typical value of 0.5. This reformulation allows for the utility of statistical

paraphrase detection models without explicit computation of semantic metric spaces.

$$S_L(s_1, s_2) = \begin{cases} 1, & P(D_L(s_1, s_2) \leq \delta_L) \geq \epsilon \\ 0, & P(D_L(s_1, s_2) \leq \delta_L) < \epsilon \end{cases} \quad (4)$$

With the aforementioned simplifications, we now re-write our equation for approximate isometry as follows:

$$S_X(s_1, s_2) = S_Y(f(s_1), f(s_2)) \quad (5)$$

3 Related work

Based on a survey of recent literature in Natural Language Processing (NLP) and NMT, we were unable to find explicitly similar studies to our research. However, we would argue that the closest field in NLP to this research would be *adversarial paraphrasing*.

Michel et al. (2019) describes adversarial paraphrasing in the purview of machine translation as constructing paraphrases that are “*meaning preserving on the source-side, but meaning-destroying on the target-side*”. For the sake of comparison, we would mildly *paraphrase* this description of adversarial paraphrasing to “*the process of perturbing an input sentence such that it is semantically equivalent on the source-side, but semantically inequivalent on the target-side*”.

In this sense, the study of adversarial paraphrasing in machine translation could be interpreted as a targetted probe into semantic *anisometry* of NMT models, compared to our research which would be an untargetted probe into semantic isometry of NMT models. Therefore adversarial paraphrasing, while having the opposite intent, is still highly similar to our research.

Michel et al. (2019): This research lays out the framework for evaluating adversarial perturbations in sequence-to-sequence models. Additionally, this research compared three automatic sequence evaluation metrics, specifically BLEU (Papineni et al., 2002), METEOR (Denkowski and Lavie, 2014) and chrF₂ (Popović, 2015), against human judgment for evaluating semantic similarity. Results from their experiments showed that chrF₂ correlates best out of the three similarity metrics with human judgment for semantic similarity detection. We utilize this finding in later parts of our study and

attempt to compare the outputs of our paraphrase detection models with respective chrF₂ scores.

Fadaee and Monz (2020): This research lays out a simple framework for constructing adversarial paraphrases through logical operations such as word insertion/deletion and numerical/gender substitution. The research correspondingly showed that such minor modifications could lead to disproportionately larger changes in translation outputs; thereby showing an adversarial effect. This research ultimately claimed that modern NMT models are generally *volatile*, or vulnerable, to targeted adversarial attacks. We attempt to compare this claim with our findings in later parts of this study.

4 Experimental setup

4.1 Data sets

Below we describe the key data sets that we were used in our research.

4.1.1 WMT19 en-de references and corresponding paraphrases

Freitag et al. (2020) builds on the premise that while automatic evaluation metrics, such as BLEU, are important for NMT model evaluation; the presence of diverse translation references is also critical. Motivated by the observation that typical references show poor diversity, Freitag et al. (2020) focuses on two goals; namely creating additional high quality WMT19 test references, as well as paraphrasing both existing (or legacy) and additional WMT19 test references in the en→de translation direction.

These services were ultimately rendered by a professional translation service using different sets of linguists for different tasks to reduce systematic bias. While these additional references serve the purpose of diversifying evaluation references, we see them as a source of high-quality semantically equivalent paraphrases with varied lexical and syntactical features. We therefore use these additional references as our input paraphrase sentences to investigate semantic isometry. Below are the key resulting data sets for en→de that were used in our research.

WMT19 legacy test references: This refers to the existing newstest2019 translation references without any additional references or paraphrasing. For abbreviation purposes, we refer to this data set as WMT.

WMT19 additional test references: This refers to additional references produced as a result of Freitag et al. (2020). For abbreviation purposes, we refer to this data set as AR.

WMT19 legacy test paraphrased references: This refers to the paraphrased version of the existing newstest2019 translation references produced as a result of Freitag et al. (2020). For abbreviation purposes, we refer to this data set as WMT.p.

WMT19 additional test paraphrased references: This refers to the paraphrased version of the additional translation references produced as a result of Freitag et al. (2020). For abbreviation purposes, we refer to this data set as AR.p.

For brevity, we further simplify the aforementioned data sets as follows:

$$\text{WMT19 Legacy} = \{\text{WMT} \cup \text{WMT.p}\} \quad (6)$$

$$\text{WMT19 AR} = \{\text{AR} \cup \text{AR.p}\} \quad (7)$$

4.1.2 PAWS-X

PAWS-X is a cross-lingual adversarial data set for paraphrase identification released by Google Research (Yang et al., 2019). PAWS-X stems originally from the PAWS data set released by Zhang et al. (2019) which is an abbreviation for Paraphrase Adversaries from Word Scrambling.

The original motivation behind the PAWS data set was that existing paraphrase detection data sets lacked non-paraphrase sentence pairs with high lexical overlap. The PAWS data set was therefore released to drive progress in creating models that utilize fine-grained structure and context of sentence pairs.

The PAWS data set contains 108,463 paraphrase and non-paraphrase sentence pairs with high lexical overlap. These sentence pairs were bulk sourced from Wikipedia and Quora Question Pairs; followed by controlled word swapping and back translation to create challenging sentence pairs for paraphrase detection. The generated sentence pairs were finally evaluated for fluency and general quality by human raters.

As noted in Yang et al. (2019), one limitation of adversarially generated data sets such as PAWS is their pre-dominant focus on the English language. In order to address this issue, Yang et al. (2019) released PAWS-X; which consists of 23,659 human translated evaluation sentence pairs and 296,406 machine-translated training sentence pairs



Figure 2: Training and validation cross entropy loss for Scaling NMT WMT16 Transformer

derived from the Wikipedia subset of the original PAWS data set. These sentence pairs were translated from English to six typologically distinct languages; namely French, Spanish, German, Chinese, Japanese and Korean.

The release of PAWS-X provides many advantages to the field of NLP, particularly the creation of a new benchmark to promote research in multilingual and zero-shot paraphrase detection. This can already be seen by the incorporation of PAWS-X into Google’s recent Cross-lingual Transfer Evaluation of Multilingual Encoders (XTREME) benchmark system (Hu et al., 2020).

4.1.3 WMT16 de-en

In this study, we replicate a non-SOTA NMT model from scratch based off the Scaling NMT WMT16 workflow (Ott et al., 2018). While the original implementation in Ott et al. (2018) is based on the en→de translation direction, our implementation trains a NMT model in the reverse translation direction; specifically de→en.

For this, we use WMT16 de→en training data with 4.5M sentence pairs, newstest2013 as our validation set and newstest2014 as our test set. We utilize a vocabulary of 32K symbols based off a joint source and target byte-pair encoding (BPE; Sennrich et al. 2015).

4.2 Models

Below we describe the key models that were used in our research.

4.2.1 FAIR WMT19 Transformer

We utilize FAIR’s winning WMT19 single Transformer model as our SOTA NMT model. Focusing particularly on the de→en translation direction, the FAIR WMT19 Transformer was the top

performing model in WMT19 with a SacreBLEU (Post, 2018) score of 40.8.

As per Ng et al. (2019), the key factors that led to SOTA performance include langid filtering of crawled bitext data, large-scale back translation as a form of data augmentation and noisy channel model reranking. We utilized this model directly from the fairseq API (Ott et al., 2019).

4.2.2 Scaling NMT WMT16 Transformer

We replicate the Scaling NMT WMT16 Transformer based on Ott et al. (2018) by training it from scratch. However, we swap the translation direction from en→de to de→en; such that we can ultimately use this model to translate WMT19 paraphrases from de→en. We intentionally choose this workflow since it would produce a non-SOTA transformer which would be useful for us downstream to introduce performance-dependent variance in the translation of WMT19 paraphrases.

Besides the aforementioned modification, we follow the same setup as per Ott et al. (2018). Specifically, we use a “big” transformer model based off Vaswani et al. (2017); with 6 blocks in the encoder and decoder networks. This model has a total of 210M parameters.

During training, we apply dropout (Srivastava et al., 2014) with probability 0.3 and utilize the Adam optimizer (Kingma and Ba, 2014) with $\beta_1 = 0.9$ and $\beta_2 = 0.98$. We use a learning rate schedule where the learning rate increases linearly for 4,000 steps from $1e-7$ until $1e-3$. The learning rate then decays proportionally to the inverse square root of the number of training steps. We utilize label smoothing with weight 0.1 for the uniform prior distribution over the vocabulary (Pereyra et al., 2017). We use large batch sizes with the

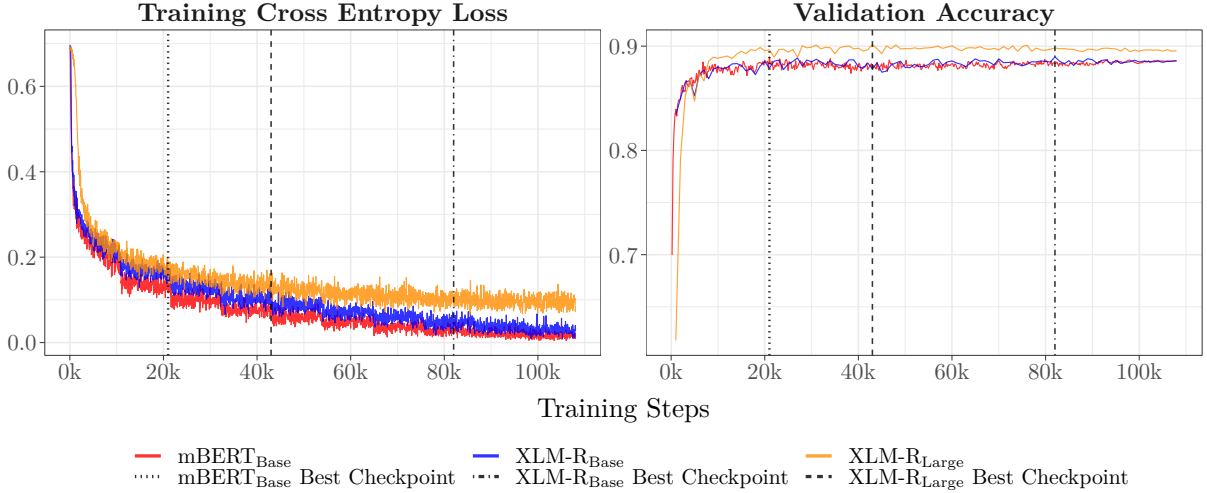


Figure 3: Training loss and validation accuracy w.r.t. training steps for paraphrase detection models

maximum number of tokens per batch being 7168. Furthermore, we apply gradient accumulation for 8 steps before updating the model; which is known as `update-freq` in the `fairseq` API. We also exploit `fairseq`’s half precision floating point (FP16) functionality for more efficient training.

Finally, we train this model for 6 days on a single NVIDIA Tesla-V100 16GB GPU. During training, we monitor the validation loss and enable checkpoint-saving for the best performing checkpoint on the validation set. We train the model up until $\sim 285K$ updates.

Our best performing checkpoint was saved at $\sim 180K$ updates as seen in Figure 2. For evaluation on the test set, we utilize beam search with a beam width of 5. Our final Scaling NMT WMT16 Transformer achieved a SacreBLEU (Post, 2018) score² of 31.0 on the `newstest2014` test set.

4.2.3 Paraphrase detection models

As noted in equation 4, paraphrase detection models are useful in computing probabilistic semantic-equivalence spaces, or otherwise the S_L function. We follow a similar framework as that detailed in Google’s XTREME benchmark (Hu et al., 2020) and fine-tune pre-trained multilingual transformer language models on the PAWS-X paraphrase detection task. We focus specifically on three multilingual transformer language models, specifically `mBERTBase` (104 languages; 172M parameters; Devlin et al. 2019), `XLM-RBase` (100 languages; 270M parameters; Conneau et al. 2019) and `XLM-RLarge`

²SacreBLEU signature:
BLEU+case.mixed+lang.de-en+numrefs.1+smooth.exp+
test.wmt14/full+tok.13a+version.1.4.12

Language	mBERT _B	XLM-R _B	XLM-R _L
en	0.940	0.946	0.960
de	0.898	0.900	0.912
es	0.908	0.922	0.928
fr	0.922	0.917	0.933
ja	0.836	0.836	0.859
ko	0.841	0.847	0.870
zh	0.854	0.861	0.876
μ	0.886	0.890	0.906

Table 1: Language-specific summary of macro- F_1 scores of paraphrase detection models on the PAWS-X test set; languages are abbreviated based on ISO 639-1; B and L refer to base and large respectively

(100 languages; 550M parameters; Conneau et al. 2019) using HuggingFace’s `transformers` library (Wolf et al., 2019) with model variants optimized for sequence classification.

While our implementation is similar to that of Google’s XTREME benchmark, we modify some aspects of the workflow to suit our needs. Most importantly, we fine-tune our multilingual language models on PAWS-X training data from all 7 languages instead of only English in order to reap the benefits of diverse multilingual data.

For all models, we enforce a maximum sequence length of 128 tokens since PAWS-X sentence pairs generally fit it into this range. We use a modified Adam optimizer with decoupled weight decay regularization (Loshchilov and Hutter, 2017) with $\beta_1 = 0.9$ and $\beta_2 = 0.999$. We train all models for 10 epochs or $\sim 110K$ updates with a global batch size of 32. We also use a linearly decaying learn-

ing rate schedule without warmup steps. Lastly, we monitor accuracy on the PAWS-X validation set for all languages in order to determine the best performing checkpoint.

Specific to mBERT_{Base} and XLM-R_{Base}, we use a batch size of 32 without gradient accumulation and an initial learning rate of 2e-5. As for XLM-R_{Large}, we use an initial learning rate of 1e-6 and local batch size of 8 with 4 gradient accumulation steps to curb GPU out-of-memory (OOM) issues.

We fine-tune mBERT_{Base}, XLM-R_{Base} and XLM-R_{Large} for 14 hours, 15 hours and 2.5 days on a single NVIDIA Geforce GTX 1080-Ti 12GB GPU respectively. The best checkpoints are achieved and saved at $\sim 20K$, $\sim 80K$ and $\sim 40K$ updates respectively, as seen in Figure 3.

As seen in Table 1, all three models perform well especially on our target languages of en and de. Overall, the best performing model on the PAWS-X test set is XLM-R_{Large} with a macro-F₁ of 0.906.

4.3 Evaluation protocols

Given our WMT19 Legacy and WMT19 AR paraphrase datasets defined in equations 6 and 7, we then translate all the aforementioned pairs of paraphrases using both the FAIR WMT19 Transformer and the Scaling NMT WMT16 Transformer. With this, we have the source paraphrases along with their target-side translations. We then use these samples for further investigation.

4.3.1 Isometry on binary semantic-equivalence spaces

Here we present our evaluation protocols for investigating the isometric properties of NMT models on binary semantic-equivalence spaces.

Vector representation: We modify our representation of the S_L relations in equation 5 in order to have a more concise vectorized form of the relationship. We assign the probability threshold ϵ from equation 4 a constant value of 0.5 for all computations of S_L in this research.

$$\mathbf{S}_{\mathbf{XY}} = \begin{bmatrix} S_X(s_1, s_2) \\ S_Y(f(s_1), f(s_2)) \end{bmatrix} \quad (8)$$

$$\mathbf{S}_{\mathbf{XY}}^T = [S_X(s_1, s_2) \quad S_Y(f(s_1), f(s_2))] \quad (9)$$

Multi-model decisions: Given this formulation of $\mathbf{S}_{\mathbf{XY}}^T$, it is worth noting that each of the three

paraphrase detection models would compute a separate $\mathbf{S}_{\mathbf{XY}}^T$ term. Since our paraphrase detection models were evaluated to have a non-zero (albeit small and similar) error rates, we decide to use the $\mathbf{S}_{\mathbf{XY}}^T$ outputs of all three models and compute the statistical mode of the three outputs. Such a majority decision would provide more confidence in any particular decision from the models. For simplicity, we assign the statistical mode function with the symbol \mathbf{M} . We therefore compute $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T)$ to check for a majority decision over the three paraphrase detection models. We assign the empty set \emptyset in case no majority decision exists.

Discrete possibilities: Since the output of S_L falls in the binary set of $\{0, 1\}$, we can effectively compute all five possibilities of the consequent majority decision $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T)$. We describe these below.

1. $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T) = [0, 0]$: Both the source and target pairs of sentences were evaluated to not be paraphrases and are therefore not semantically equivalent on both sides. This qualifies as isometric behaviour.
2. $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T) = [1, 1]$: Both the source and target pairs of sentences were evaluated to be paraphrases and are therefore semantically equivalent on both sides. This qualifies as isometric behaviour.
3. $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T) = [0, 1]$: The source sentences are evaluated to not be paraphrases while the target sentences are evaluated to be paraphrases. This implies that translation resulted in the sentence pair becoming semantically equivalent while they were not before. This qualifies as anisometric behaviour and could be an interesting scenario to investigate further.
4. $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T) = [1, 0]$: The source sentences are evaluated to be paraphrases while the target sentences are evaluated to not be paraphrases. This implies that translation resulted in the sentence pair becoming semantically inequivalent while they were not before. This qualifies as anisometric behaviour and could imply weak performance of the model on one of the sentences; or possibly some adversarial paraphrasing.
5. $\mathbf{M}(\mathbf{S}_{\mathbf{XY}}^T) = \emptyset$: This implies that all models had different decisions and therefore no majority decision was reached.




Figure 4: Normalized contour densities for target paraphrase softmax score against source paraphrase softmax score; grouped by NMT models (top), paraphrase detection models (top) and input data sets (right)

Frequency analysis: Given the five aforementioned possibilities of $M(S_{XY}^T)$, we measure the frequency of each possibility in each model’s input and output sentences. This give us an insight into the isometric behaviour of each model and would also allow us to compare both models.

Sentence-level analysis: As mentioned before, some anisometric possibilities of $M(S_{XY}^T)$, such as $[0, 1]$ and $[1, 0]$, could prove interesting for further investigation. For such cases, we probe further and analyze sentences individually.

4.3.2 Relationship between chrF_2 and semantic-equivalence

We return to one of the observations from Michel et al. (2019), specifically that the chrF_2 automatic evaluation metric (Popović, 2015) tends to correlate most positively with human judgment of semantic similarity compared to BLEU (Papineni et al., 2002) and METEOR (Denkowski and Lavie, 2014). We attempt to use our investigation of isometry on semantic-equivalence spaces to further investigate the relationship between semantic equality and the chrF_2 score. We replicate the chrF_2 setup from Michel et al. (2019) by using the default sacrebleu implementation (Post, 2018) with a n -gram upper limit of 6 and β value of 2.

Non-commutativity of chrF_2 : While replicating the chrF_2 setup of Michel et al. (2019), we observe that this particular formulation with $\beta = 2$ results in the chrF_2 metric being non-commutative:

$$\text{chrF}_2(s_1, s_2) \neq \text{chrF}_2(s_2, s_1) \quad (10)$$

Non-commutativity is an emergent property of chrF_2 with $\beta = 2$ since this would assign two times more weight to recall than precision (Popović, 2015); which places an internal bias on input order. While this non-commutative formulation of chrF_2 would be useful for evaluating NMT models with explicit hypotheses and references, this would not be optimal as a semantic similarity metric since one would expect a semantic similarity metric to be commutative and unbiased towards input order.

Commutative variant of chrF_2 : As an alternative, we simulate commutativity by averaging the chrF_2 values for both input orders and introduce $\overline{\text{chrF}_2}$ as a commutative variant for chrF_2 .

$$\overline{\text{chrF}_2}(s_1, s_2) = \frac{\text{chrF}_2(s_1, s_2) + \text{chrF}_2(s_2, s_1)}{2} \quad (11)$$

$$\therefore \overline{\text{chrF}_2}(s_1, s_2) = \overline{\text{chrF}_2}(s_2, s_1) \quad (12)$$

We see this as a more optimal alternative than changing β to 1 since this might veer further away from the experimental setup of Michel et al. (2019). Such commutative variants of automatic similarity metrics have also been considered for BLEU and METEOR in Wieting et al. (2019) and were termed *symmetric* instead of commutative.

Mapping $\overline{\text{chrF}_2}$ to S_L : In order to investigate the relationship between $\overline{\text{chrF}_2}$ and semantic-equivalence (or S_L), we start by assigning sentence pairs with respective paired chrF_2 and S_L values. In order to do this, we refer back to the majority

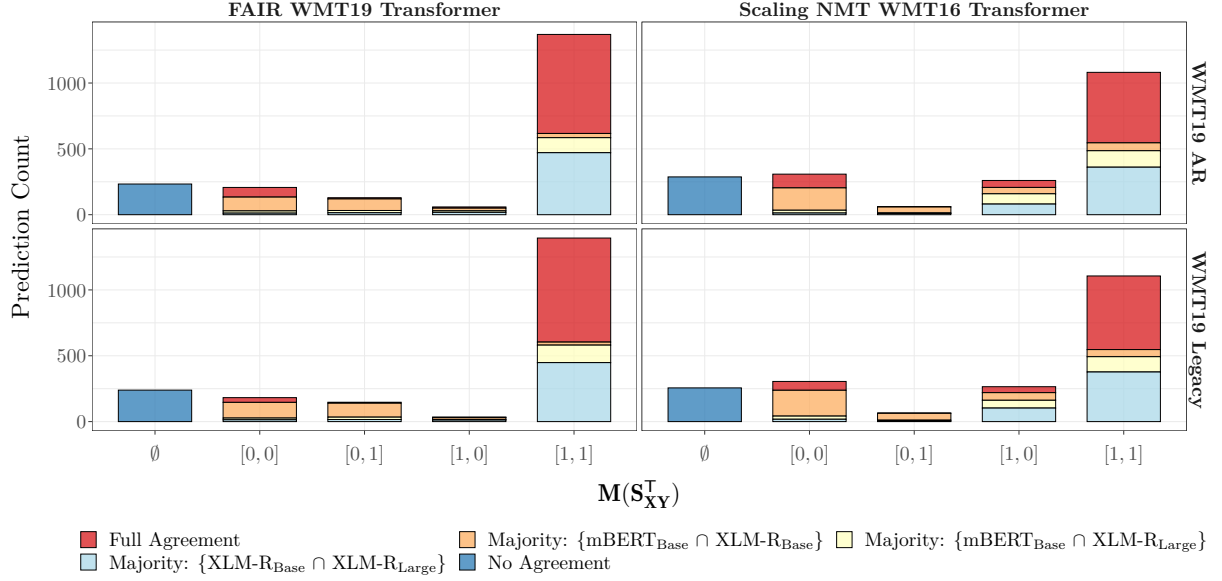


Figure 5: Overview of $M(S_{XY}^T)$ decisions and frequencies w.r.t. NMT models (top) and input data sets (right); filling colors indicate finer details on the type of majority decision

decisions of the paraphrase detection models and remove all sentence pairs where $M(S_{XY}^T) = \emptyset$.

We assume all remaining sentence pairs have been *confidently* tagged by the paraphrase detection models and we correspondingly assign each sentence pair with respective S_L and $\overline{chrF_2}$ values for source and target sides.

Correlation between $\overline{chrF_2}$ and S_L : Finally, we replicate a similar statistical procedure as per Michel et al. (2019) and compute the Pearson correlation coefficient r and the corresponding statistical significance t -test. For the t -test, we set the null hypothesis H_0 to be that there exists a non-positive correlation between $\overline{chrF_2}$ and S_L ; while the alternative hypothesis H_1 is that there exists a positive correlation between $\overline{chrF_2}$ and S_L .

4.4 Results

References

- Daniel Cer, Mona Diab, Eneko Agirre, Inigo Lopez-Gazpio, and Lucia Specia. 2017. Semeval-2017 task 1: Semantic textual similarity-multilingual and cross-lingual focused evaluation. *arXiv preprint arXiv:1708.00055*.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Unsupervised cross-lingual representation learning at scale. *arXiv preprint arXiv:1911.02116*.
- Harold Scott Macdonald Coxeter. 1961. Introduction to geometry.
- Michael Denkowski and Alon Lavie. 2014. Meteor universal: Language specific translation evaluation for any target language. In *Proceedings of the ninth workshop on statistical machine translation*, pages 376–380.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Marzieh Fadaee and Christof Monz. 2020. The unreasonable volatility of neural machine translation models. *arXiv preprint arXiv:2005.12398*.
- Markus Freitag, David Grangier, and Isaac Caswell. 2020. Bleu might be guilty but references are not innocent. *ArXiv*, abs/2004.06063.
- C. Hegde, A. C. Sankaranarayanan, W. Yin, and R. G. Baraniuk. 2015. Numax: A convex approach for learning near-isometric linear embeddings. *IEEE Transactions on Signal Processing*, 63(22):6109–6121.
- Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. 2020. [Xtreme: A massively multilingual multi-task benchmark for evaluating cross-lingual generalization](#). *CoRR*, abs/2003.11080.
- Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- Ilya Loshchilov and Frank Hutter. 2017. [Fixing weight decay regularization in adam](#). *CoRR*, abs/1711.05101.
- Paul Michel, Xian Li, Graham Neubig, and Juan Miguel Pino. 2019. On evaluation of adversarial perturbations for sequence-to-sequence models. *arXiv preprint arXiv:1903.06620*.
- Nathan Ng, Kyra Yee, Alexei Baevski, Myle Ott, Michael Auli, and Sergey Edunov. 2019. Facebook fair’s wmt19 news translation task submission. *arXiv preprint arXiv:1907.06616*.
- Myle Ott, Sergey Edunov, Alexei Baevski, Angela Fan, Sam Gross, Nathan Ng, David Grangier, and Michael Auli. 2019. fairseq: A fast, extensible toolkit for sequence modeling. In *Proceedings of NAACL-HLT 2019: Demonstrations*.
- Myle Ott, Sergey Edunov, David Grangier, and Michael Auli. 2018. Scaling neural machine translation. In *Proceedings of the Third Conference on Machine Translation (WMT)*.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pages 311–318.
- Gabriel Pereyra, George Tucker, Jan Chorowski, Łukasz Kaiser, and Geoffrey Hinton. 2017. Regularizing neural networks by penalizing confident output distributions. *arXiv preprint arXiv:1701.06548*.
- Maja Popović. 2015. chrF: character n-gram f-score for automatic mt evaluation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395.
- Matt Post. 2018. [A call for clarity in reporting BLEU scores](#). In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186–191, Belgium, Brussels. Association for Computational Linguistics.
- Rico Sennrich, Barry Haddow, and Alexandra Birch. 2015. Neural machine translation of rare words with subword units. *arXiv preprint arXiv:1508.07909*.
- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. 2014. Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1):1929–1958.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in neural information processing systems*, pages 5998–6008.
- John Wieting, Taylor Berg-Kirkpatrick, Kevin Gimpel, and Graham Neubig. 2019. [Beyond BLEU: training neural machine translation with semantic similarity](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4344–4355, Florence, Italy. Association for Computational Linguistics.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, R’emi Louf, Morgan Funtowicz, and Jamie Brew. 2019. Huggingface’s trans-

formers: State-of-the-art natural language processing. *ArXiv*, abs/1910.03771.

Yinfei Yang, Yuan Zhang, Chris Tar, and Jason Baldridge. 2019. Paws-x: A cross-lingual adversarial dataset for paraphrase identification. In *Proc. of EMNLP*.

Yuan Zhang, Jason Baldridge, and Luheng He. 2019. Paws: Paraphrase adversaries from word scrambling. *arXiv preprint arXiv:1904.01130*.